

Tidy Time Series & Forecasting in R

10. Forecast reconciliation

bit.ly/fable2020



Outline

- 1 Hierarchical and grouped time series
- 2 Forecast reconciliation
- 3 Example: Australian tourism
- 4 Lab Session 20

Outline

- 1 Hierarchical and grouped time series
- 2 Forecast reconciliation
- 3 Example: Australian tourism
- 4 Lab Session 20

Australian Pharmaceutical Benefits Scheme



PBS sales

PBS

```
## # A tibble: 65,219 x 9 [1M]
## # Key:      Concession, Type, ATC1, ATC2 [336]
##       Month Concession Type ATC1 ATC1_desc ATC2 ATC2_desc Scripts
##       <mt> <chr>      <chr> <chr> <chr>      <chr> <chr>      <dbl>
##  1 1991 Jul Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 18228
##  2 1991 Aug Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 15327
##  3 1991 Sep Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 14775
##  4 1991 Oct Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 15380
##  5 1991 Nov Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 14371
##  6 1991 Dec Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 15028
##  7 1992 Jan Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 11040
##  8 1992 Feb Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 15165
##  9 1992 Mar Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 16898
## 10 1992 Apr Concessio~ Co-p~ A      Alimenta~ A01    STOMATOL~ 18141
## # ... with 65,209 more rows, and 1 more variable: Cost <dbl>
```

ATC drug classification

- A Alimentary tract and metabolism
- B Blood and blood forming organs
- C Cardiovascular system
- D Dermatologicals
- G Genito-urinary system and sex hormones
- H Systemic hormonal preparations, excluding sex hormones and insulins
- J Anti-infectives for systemic use
- L Antineoplastic and immunomodulating agents
- M Musculo-skeletal system
- N Nervous system
- P Antiparasitic products, insecticides and repellents
- R Respiratory system
- S Sensory organs
- V Various

ATC drug classification

ATC1: 14 classes

A

Alimentary tract and metabolic

ATC2: 84 classes

A10

Drugs used in diabetes

A10B

Blood glucose lowering drugs

A10BA


Biguanides

A10BA02

Metformin

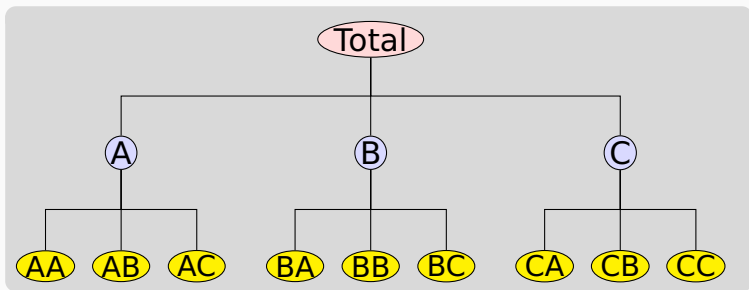
Hierarchical time series

A **hierarchical time series** is a collection of several time series that are linked together in a hierarchical structure.



Hierarchical time series

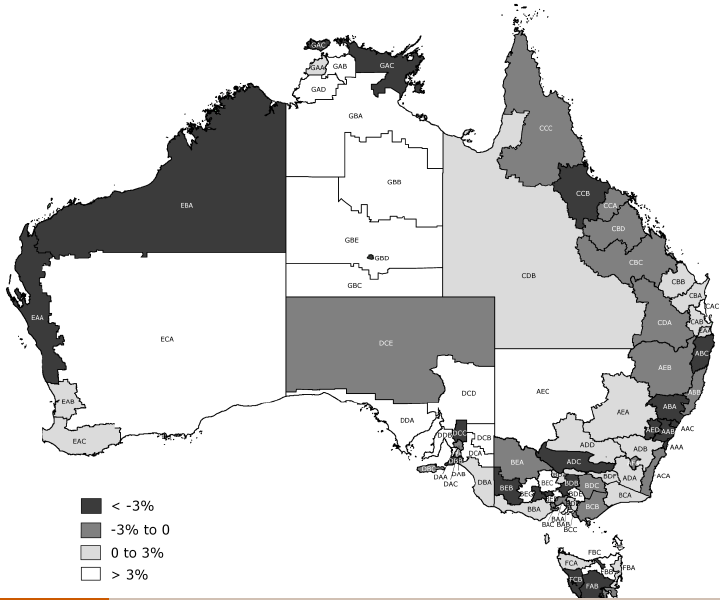
A **hierarchical time series** is a collection of several time series that are linked together in a hierarchical structure.



Examples

- PBS sales by ATC groups
- Tourism demand by states, zones, regions

Australian tourism



Australian tourism

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
```

```
## # Key:           Region, State, Purpose [304]
```


```
##   Quarter Region   State           Purpose   Trips
##   <qtr> <chr>      <chr>           <chr>      <dbl>
## 1 1998 Q1 Adelaide South Australia Business   135.
## 2 1998 Q2 Adelaide South Australia Business   110.
## 3 1998 Q3 Adelaide South Australia Business   166.
## 4 1998 Q4 Adelaide South Australia Business   127.
## 5 1999 Q1 Adelaide South Australia Business   137.
## 6 1999 Q2 Adelaide South Australia Business   200.
## 7 1999 Q3 Adelaide South Australia Business   169.
## 8 1999 Q4 Adelaide South Australia Business   134.
## 9 2000 Q1 Adelaide South Australia Business   154.
## 10 2000 Q2 Adelaide South Australia Business   169.
## # ... with 24,310 more rows
```

Australian tourism

- Quarterly data on visitor night from 1998:Q1 – 2013:Q4
- From: *National Visitor Survey*, based on annual interviews of 120,000 Australians aged 15+, collected by Tourism Research Australia.
- Split by 7 states, 27 zones and 76 regions (a geographical hierarchy)
- Also split by purpose of travel
 - ▶ Holiday
 - ▶ Visiting friends and relatives (VFR)
 - ▶ Business
 - ▶ Other
- 304 bottom-level series


Grouped time series

A **grouped time series** is a collection of time series that can be grouped together in a number of non-hierarchical ways.



Grouped time series

A **grouped time series** is a collection of time series that can be grouped together in a number of non-hierarchical ways.



Examples

- Tourism by state and purpose of travel
- Retail sales by product groups/sub groups, and by countries/regions

Creating aggregates

```
PBS %>%  
  aggregate_key(ATC1 / ATC2, Scripts = sum(Scripts)) %>%  
  filter(Month == yearmonth("1991 Jul")) %>%  
  print(n = 18)
```

```
## # A tibble: 98 x 4 [1M]  
## # Key:      ATC1, ATC2 [98]  
##      Month ATC1      ATC2      Scripts  
##      <mtm> <chr>      <chr>      <dbl>  
## 1 1991 Jul <aggregated> <aggregated> 8090395  
## 2 1991 Jul A      <aggregated> 799025  
## 3 1991 Jul B      <aggregated> 109227  
## 4 1991 Jul C      <aggregated> 1794995  
## 5 1991 Jul D      <aggregated> 299779  
## 6 1991 Jul G      <aggregated> 300931  
## 7 1991 Jul H      <aggregated> 112114  
## 8 1991 Jul J      <aggregated> 1151681  
## 9 1991 Jul L      <aggregated> 24580  
## 10 1991 Jul M     <aggregated> 562956  
## 11 1991 Jul N     <aggregated> 1546023  
## 12 1991 Jul P     <aggregated> 47661  
## 13 1991 Jul R     <aggregated> 859273  
## 14 1991 Jul S     <aggregated> 391639  
## 15 1991 Jul V     <aggregated> 38705  
## 16 1991 Jul Z     <aggregated> 51806  
## 17 1991 Jul A     A01          22615  
## 18 1991 Jul A     A02          299251  
## # ... with 80 more rows
```

Creating aggregates

```
tourism %>%  
  aggregate_key(Purpose * (State / Region), Trips = sum(Trips)) %>%  
  filter(Quarter == yearquarter("1998 Q1")) %>%  
  print(n = 15)
```

```
## # A tibble: 425 x 5 [1Q]  
## # Key:      Purpose, State, Region [425]  
##   Quarter Purpose      State      Region      Trips  
##   <qtr> <chr>      <chr>      <chr>      <dbl>  
## 1 1998 Q1 <aggregated> <aggregated> <aggregated> 23182.  
## 2 1998 Q1 Business <aggregated> <aggregated> 3599.  
## 3 1998 Q1 Holiday <aggregated> <aggregated> 11806.  
## 4 1998 Q1 Other <aggregated> <aggregated> 680.  
## 5 1998 Q1 Visiting <aggregated> <aggregated> 7098.  
## 6 1998 Q1 <aggregated> ACT ~ <aggregated> 551.  
## 7 1998 Q1 <aggregated> New South Wales~ <aggregated> 8040.  
## 8 1998 Q1 <aggregated> Northern Territ~ <aggregated> 181.  
## 9 1998 Q1 <aggregated> Queensland ~ <aggregated> 4041.  
## 10 1998 Q1 <aggregated> South Australia~ <aggregated> 1735.  
## 11 1998 Q1 <aggregated> Tasmania ~ <aggregated> 982.  
## 12 1998 Q1 <aggregated> Victoria ~ <aggregated> 6010.  
## 13 1998 Q1 <aggregated> Western Austral~ <aggregated> 1641.  
## 14 1998 Q1 <aggregated> ACT ~ Canberra ~ 551.  
## 15 1998 Q1 <aggregated> New South Wales~ Blue Mounta~ 196.  
## # ... with 410 more rows
```


Creating aggregates

- Similar to summarise() but using the key structure
- A grouped structure is specified using grp1 * grp2
- A nested structure is specified via parent / child.
- Groups and nesting can be mixed:

```
(country/region/city) * (brand/product)
```

- All possible aggregates are produced.
- These are useful when forecasting at different levels of aggregation.

Outline

- 1 Hierarchical and grouped time series
- 2 Forecast reconciliation
- 3 Example: Australian tourism
- 4 Lab Session 20

The problem

- 1 How to forecast time series at all nodes such that the forecasts add up in the same way as the original data?
- 2 Can we exploit relationships between the series to improve the forecasts?

The problem

- 1 How to forecast time series at all nodes such that the forecasts add up in the same way as the original data?
- 2 Can we exploit relationships between the series to improve the forecasts?

The solution

- 1 Forecast all series at all levels of aggregation using an automatic forecasting algorithm. (e.g., ETS, ARIMA, ...)
- 2 Reconcile the resulting forecasts so they add up correctly using least squares optimization (i.e., find closest reconciled forecasts to the original forecasts).
- 3 This is available using `reconcile()`.

Forecast reconciliation

```
tourism %>%  
  aggregate_key(Purpose * (State / Region), Trips = sum(Trips)) %>%  
  model(ets = ETS(Trips)) %>%  
  reconcile(ets_adjusted = min_trace(ets)) %>%  
  forecast(h = 2)
```

```
## # A tibble: 1,700 x 7 [1Q]  
## # Key:      Purpose, State, Region, .model [850]  
##   Purpose      State      Region      .model Quarter  
##   <chr>        <chr>        <chr>        <chr>    <qtr>  
## 1 Business    <aggregat~    <aggregat~    ets      2018 Q1  
## 2 Business    <aggregat~    <aggregat~    ets      2018 Q2  
## 3 Business    <aggregat~    <aggregat~    ets_a~   2018 Q1  
## 4 Business    <aggregat~    <aggregat~    ets_a~   2018 Q2  
## 5 Business    South Aus~    <aggregat~    ets      2018 Q1  
## 6 Business    South Aus~    <aggregat~    ets      2018 Q2  
## 7 Business    South Aus~    <aggregat~    ets_a~   2018 Q1  
## 8 Business    South Aus~    <aggregat~    ets_a~   2018 Q2  
## 9 Business    Northern ~    <aggregat~    ets      2018 Q1  
## 10 Business   Northern ~    <aggregat~    ets      2018 Q2  
## # ... with 1,690 more rows, and 2 more variables:
```

Hierarchical and grouped time series


Every collection of time series with aggregation constraints can be written as

$$\mathbf{y}_t = \mathbf{S}\mathbf{b}_t$$


where

- \mathbf{y}_t is a vector of all series at time t
- \mathbf{b}_t is a vector of the most disaggregated series at time t
- \mathbf{S} is a “summing matrix” containing the aggregation constraints.

Hierarchical time series



Hierarchical time series




y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

\mathbf{b}_t : vector of all series at bottom level in time t .

Hierarchical time series




y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

\mathbf{b}_t : vector of all series at bottom level in time t .

$$\mathbf{y}_t = \begin{pmatrix} y_t \\ y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix} = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix}$$

Hierarchical time series



y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

\mathbf{b}_t : vector of all series at bottom level in time t .

$$\mathbf{y}_t = \begin{pmatrix} y_t \\ y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix} = \underbrace{\begin{pmatrix} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}}_{\mathbf{S}} \underbrace{\begin{pmatrix} y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix}}_{\mathbf{b}_t}$$

$$\mathbf{y}_t = \mathbf{S}\mathbf{b}_t$$

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t . (In general, they will not “add up”.)

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t . (In general, they will not “add up”.)

Reconciled forecasts must be of the form:

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}\mathbf{G}\hat{\mathbf{y}}_n(h)$$

for some matrix \mathbf{G} .

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t . (In general, they will not “add up”.)

Reconciled forecasts must be of the form:

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}\mathbf{G}\hat{\mathbf{y}}_n(h)$$

for some matrix \mathbf{G} .

- \mathbf{G} extracts and combines base forecasts $\hat{\mathbf{y}}_n(h)$ to get bottom-level forecasts.
- \mathbf{S} adds them up

Optimal combination forecasts

Main result

The best (minimum sum of variances) unbiased forecasts are obtained when $\mathbf{G} = (\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}$, where Σ_h is the h -step base forecast error covariance matrix.

Optimal combination forecasts

Main result

The best (minimum sum of variances) unbiased forecasts are obtained when $\mathbf{G} = (\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}$, where Σ_h is the h -step base forecast error covariance matrix.

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}(\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}\hat{\mathbf{y}}_n(h)$$

Problem: Σ_h hard to estimate, especially for $h > 1$.

Solutions:

- Ignore Σ_h (OLS) [`min_trace(method='ols')`]
- Assume $\Sigma_h = k_h \Sigma_1$ is diagonal (WLS) [`min_trace(method='wls')`]
- Assume $\Sigma_h = k_h \Sigma_1$ and estimate it (GLS)

Features

- Covariates can be included in initial forecasts.
- Adjustments can be made to initial forecasts at any level.
- Very simple and flexible method. Can work with *any* hierarchical or grouped time series.
- Conceptually easy to implement: regression of base forecasts on structure matrix.

Outline


- 1 Hierarchical and grouped time series
- 2 Forecast reconciliation
- 3 Example: Australian tourism
- 4 Lab Session 20

Example: Australian tourism

```
tourism_agg <- tourism %>%  
  aggregate_key(Purpose * (State / Region),  
    Trips = sum(Trips)  
  )  
fc <- tourism_agg %>%  
  filter_index(. ~ "2015 Q4") %>%  
  model(ets = ETS(Trips)) %>%  
  reconcile(ets_adjusted = min_trace(ets)) %>%  
  forecast(h = "2 years")
```


Example: Australian tourism

```
fc %>%  
  filter(is_aggregated(Purpose) & is_aggregated(State)) %>%  
  autoplot(tourism_agg, level = 95)
```




Example: Australian tourism

```
fc %>%  
  filter(is_aggregated(Purpose) & State == "Victoria" &  
         is_aggregated(Region)) %>%  
  autoplot(tourism_agg, level = 95)
```



Example: Australian tourism

```
fc %>%  
  filter(is_aggregated(Purpose) & Region == "Melbourne") %>%  
  autoplot(tourism_agg, level = 95)
```




Example: Australian tourism

```
fc %>%
```


```
  filter(is_aggregated(Purpose) & Region == "Snowy Mountains") %>%
```

```
  autoplot(tourism_agg, level = 95)
```



Example: Australian tourism


```
fc %>%  
  filter(Purpose == "Holiday" & Region == "Barossa") %>%  
  autoplot(tourism_agg, level = 95)
```



Example: Australian tourism

```
fc %>%
```

```
  filter(is_aggregated(Purpose) & Region == "MacDonnell") %>%  
  autoplot(tourism_agg, level = 95)
```



Example: Australian tourism

```
fc <- tourism_agg %>%  
  filter_index(. ~ "2015 Q4") %>%  
  model(  
    ets = ETS(Trips),  
    arima = ARIMA(Trips)  
  ) %>%  
  mutate(  
    comb = (ets + arima) / 2  
  ) %>%  
  reconcile(  
    ets_adj = min_trace(ets),  
    arima_adj = min_trace(arima),  
    comb_adj = min_trace(comb)  
  ) %>%  
  forecast(h = "2 years")
```

Forecast evaluation

```
fc %>% accuracy(tourism_agg)
```

```
## # A tibble: 2,550 x 12
```

```
##   .model Purpose      State      Region      .type      ME
##   <chr>  <chr>      <chr>      <chr>      <chr>      <dbl>
## 1 arima  Business    <aggregat~ <aggregat~ Test      685.
## 2 arima  Business    South Aus~ <aggregat~ Test       49.9
## 3 arima  Business    Northern ~ <aggregat~ Test       22.2
## 4 arima  Business    Western A~ <aggregat~ Test     -138.
## 5 arima  Business    Victoria ~ <aggregat~ Test      232.
## 6 arima  Business    New South~ <aggregat~ Test      153.
## 7 arima  Business    Queenslan~ <aggregat~ Test      81.8
## 8 arima  Business    ACT        ~ <aggregat~ Test      35.9
## 9 arima  Business    Tasmania ~ <aggregat~ Test      28.8
## 10 arima Business    South Aus~ Adelaide   Test      20.8
## # ... with 2,540 more rows, and 6 more variables:
## #   RMSE <dbl>, MAE <dbl>, MPE <dbl>, MAPE <dbl>,
## #   MASE <dbl>, ACF1 <dbl>
```

Forecast evaluation

```
fc %>%  
  accuracy(tourism_agg) %>%  
  group_by(.model) %>%  
  summarise(MASE = mean(MASE)) %>%  
  arrange(MASE)
```

```
## # A tibble: 6 x 2  
##   .model      MASE  
##   <chr>      <dbl>  
## 1 ets_adj    1.02  
## 2 comb_adj   1.02  
## 3 ets        1.04  
## 4 comb       1.04  
## 5 arima_adj  1.07  
## 6 arima      1.09
```

Outline

- 1 Hierarchical and grouped time series
- 2 Forecast reconciliation
- 3 Example: Australian tourism
- 4 Lab Session 20

Lab Session 20

- Prepare aggregations of the PBS data by Concession, Type, and ATC1.
- Use forecast reconciliation with the PBS data, using ETS, ARIMA and SNAIVE models, applied to all but the last 3 years of data.
- Which type of model works best?
- Does the reconciliation improve the forecast accuracy?
- Why doesn't the reconciliation make any difference to the SNAIVE forecasts?

rstd.io/ws-survey